

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
data = pd.read_csv("churn_data_te.csv")
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 64374 entries, 0 to 64373
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   CustomerID            64374 non-null  int64
 1   Age                   64374 non-null  int64
 2   Gender                64374 non-null  object
 3   Tenure                64374 non-null  int64
 4   Usage Frequency       64374 non-null  int64
 5   Support Calls         64374 non-null  int64
 6   Payment Delay         64374 non-null  int64
 7   Subscription Type     64374 non-null  object
 8   Contract Length       64374 non-null  object
 9   Total Spend           64374 non-null  int64
10   Last Interaction      64374 non-null  int64
11   Churn                 64374 non-null  int64
dtypes: int64(9), object(3)
memory usage: 5.9+ MB
```

```
data.head()
```

	CustomerID	Age	Gender	Tenure	Usage Frequency	Support Calls	Payment Delay	Subscription Type	Contract Length	Total Spend	Last Interaction	Churn
0	1	22	Female	25	14	4	27	Basic	Monthly	598	9	1
1	2	41	Female	28	28	7	13	Standard	Monthly	584	20	0
2	3	47	Male	27	10	2	29	Premium	Annual	757	21	0
3	4	35	Male	9	12	5	17	Premium	Quarterly	232	18	0

```
data.tail()
```

	CustomerID	Age	Gender	Tenure	Usage Frequency	Support Calls	Payment Delay	Subscription Type	Contract Length	Total Spend	Last Interaction	Churn
64369	64370	45	Female	33	12	6	21	Basic	Quarterly	947	14	1
64370	64371	37	Male	6	1	5	22	Standard	Annual	923	9	1
64371	64372	25	Male	39	14	8	30	Premium	Monthly	327	20	1
64372	64373	50	Female	18	19	7	22	Standard	Monthly	540	13	1

```
data.isnull().sum()
```

```
CustomerID      0
Age              0
Gender           0
Tenure          0
Usage Frequency  0
Support Calls    0
Payment Delay    0
Subscription Type 0
Contract Length  0
Total Spend      0
Last Interaction  0
Churn            0
dtype: int64
```

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
```

```
# Detect object columns (categorical)
categorical_cols = data.select_dtypes(include='object').columns
print("Columns that need encoding:", categorical_cols.tolist())
```

Columns that need encoding: []

```
label_encoders = {}
for column in data.select_dtypes(include='object').columns:
    le = LabelEncoder()
    data[column] = le.fit_transform(data[column])
    label_encoders[column] = le
```

```
X = data.drop('Churn', axis=1)
y = data['Churn']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
model = LogisticRegression()
model.fit(X_train, y_train)
```

LogisticRegression

```
y_pred = model.predict(X_test)
```

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

Accuracy: 0.8653980582524272

```
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

Confusion Matrix:

```
[[5929  864]
 [ 869 5213]]
```

```
print("Classification Report:\n", classification_report(y_test, y_pred))
```


Classification Report:

	precision	recall	f1-score	support
0	0.87	0.87	0.87	6793
1	0.86	0.86	0.86	6082
accuracy			0.87	12875
macro avg	0.86	0.86	0.86	12875
weighted avg	0.87	0.87	0.87	12875


```
import seaborn as sns
```

```
sns.set(style="whitegrid")
```


```
plt.figure(figsize=(6, 4))
sns.countplot(data=data, x='Churn')
plt.title('Customer Churn Count')
plt.xlabel('Churn')
plt.ylabel('Count')
plt.show()
```




```
plt.figure(figsize=(6, 4))
sns.countplot(data=data, x='Gender', hue='Churn')
plt.title('Churn by Gender')
plt.show()
```



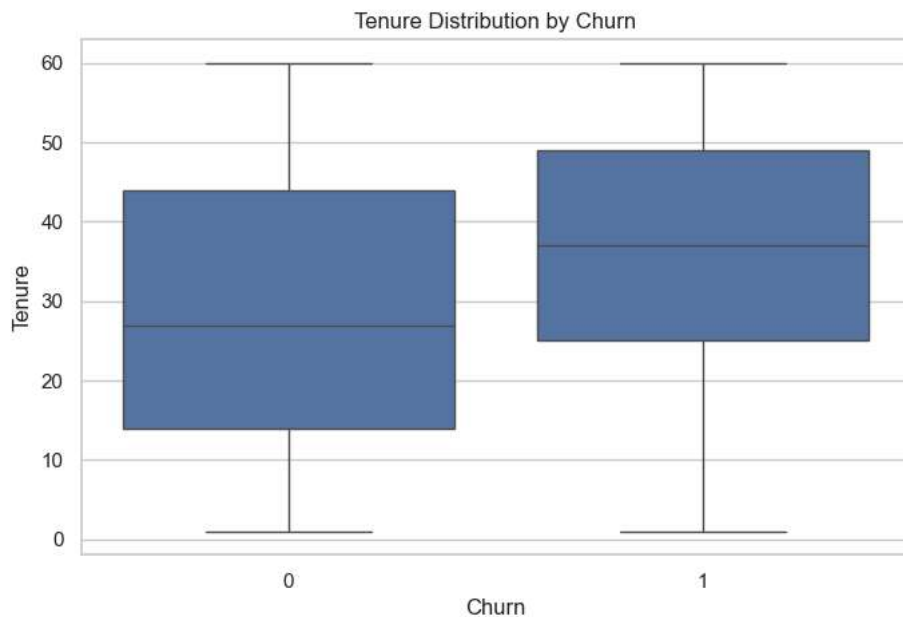
```
plt.figure(figsize=(8, 5))
sns.histplot(data=data, x='Total Spend', hue='Churn', kde=True, bins=30)
plt.title('Monthly Charges Distribution by Churn')
plt.show()
```




```
plt.figure(figsize=(8, 5))
sns.histplot(data=data, x='Last Interaction', hue='Churn', kde=True, bins=30)
plt.title('Monthly Charges Distribution by Churn')
plt.show()
```



```
plt.figure(figsize=(8, 5))
sns.boxplot(data=data, x='Churn', y='Tenure')
plt.title('Tenure Distribution by Churn')
plt.show()
```



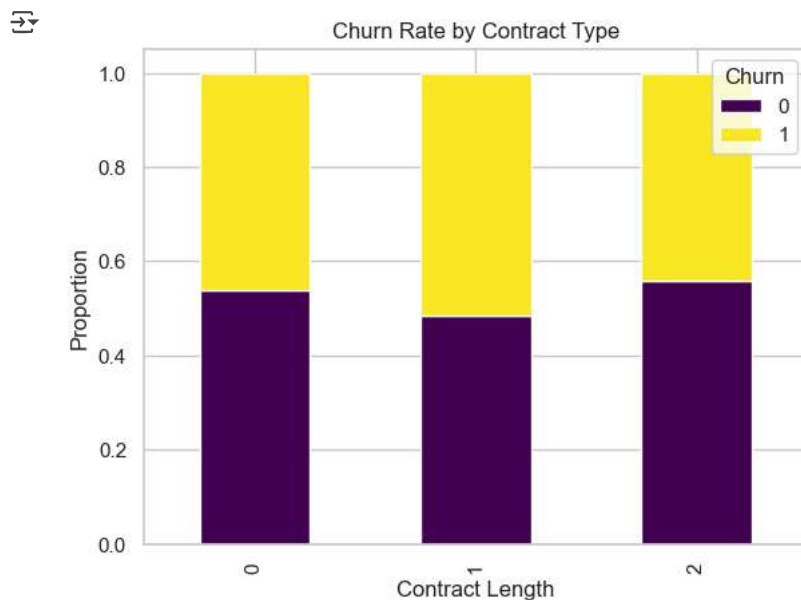
```
plt.figure(figsize=(10, 8))
corr = data.select_dtypes(include=['float64', 'int64']).corr()
sns.heatmap(corr, annot=True, fmt=".2f", cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
```



```

contract_churn = pd.crosstab(data['Contract Length'], data['Churn'], normalize='index')
contract_churn.plot(kind='bar', stacked=True, colormap='viridis')
plt.title('Churn Rate by Contract Type')
plt.ylabel('Proportion')
plt.show()

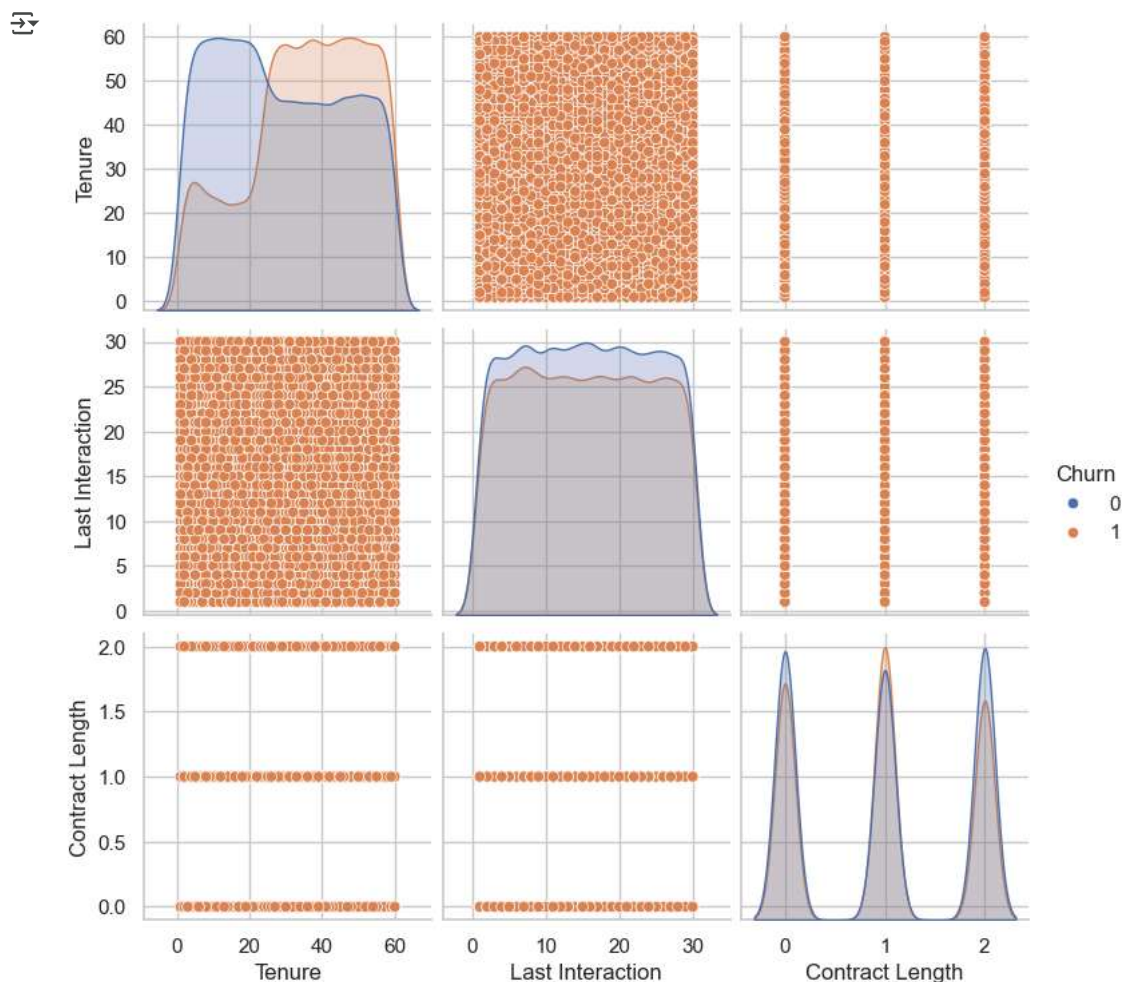
```




```

sns.pairplot(data[['Tenure', 'Last Interaction', 'Contract Length', 'Churn']], hue='Churn', diag_kind='kde')
plt.show()

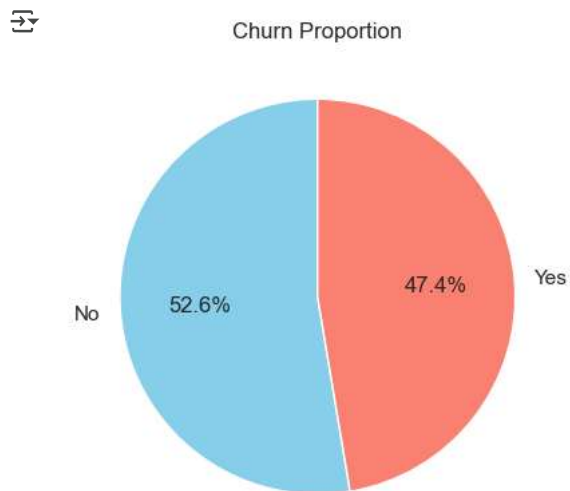
```



```
sns.violinplot(x='Churn', y='Payment Delay', data=data)
plt.title('Monthly Charges Distribution by Churn')
plt.show()
```



```
data['Churn'].value_counts().plot.pie(autopct='%1.1f%%', startangle=90, labels=['No', 'Yes'], colors=['skyblue', 'salmon'])
plt.title('Churn Proportion')
plt.ylabel('')
plt.show()
```



```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix
```

```
data.dropna(inplace=True)
X = data.drop('Churn', axis=1)
y = data['Churn']
```

```
X = pd.get_dummies(X, drop_first=True)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(X_train, y_train)
```

RandomForestClassifier

RandomForestClassifier(random_state=42)

```
y_pred = model.predict(X_test)
print(classification_report(y_test, y_pred))
```

```
precision    recall  f1-score   support

0           1.00      1.00      1.00      6793
1           1.00      1.00      1.00      6082

accuracy          1.00      1.00      1.00     12875
macro avg          1.00      1.00      1.00     12875
weighted avg          1.00      1.00      1.00     12875
```

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
# Predict on the test set
y_pred = model.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.9979029126213592

```
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", conf_matrix)
```

Confusion Matrix:

```
[[6793  0]
 [ 27 6055]]
```

```
print("Classification Report:\n", classification_report(y_test, y_pred))
```

```
Classification Report:
precision    recall  f1-score   support

0           1.00      1.00      1.00      6793
1           1.00      1.00      1.00      6082

accuracy          1.00      1.00      1.00     12875
macro avg          1.00      1.00      1.00     12875
weighted avg          1.00      1.00      1.00     12875
```

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from xgboost import XGBClassifier
```

```
X = data.drop("Churn", axis=1)
y = data["Churn"]
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```


Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.


```
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
svm_model = SVC()
```

```
knn_model = KNeighborsClassifier()
```

```
xgb_model = XGBClassifier(use_label_encoder=False, eval_metric='logloss', verbosity=0)
```

```
svm_model.fit(X_train_scaled, y_train)
```

 SVC ⓘ ?


```
SVC()
```

```
knn_model.fit(X_train_scaled, y_train)
```

 KNeighborsClassifier ⓘ ?

```
KNeighborsClassifier()
```

```
xgb_model.fit(X_train, y_train)
```

 XGBClassifier ⓘ ?

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric='logloss',
               feature_types=None, feature_weights=None, gamma=None,
               grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=None, max_leaves=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,
               multi_strategy=None, n_estimators=None, n_jobs=None,
```


```
svm_preds = svm_model.predict(X_test_scaled)
```

```
knn_preds = knn_model.predict(X_test_scaled)
```

```
xgb_preds = xgb_model.predict(X_test)
```

```
print("SVM Accuracy:", accuracy_score(y_test, svm_preds))
```

```
print("SVM Report:\n", classification_report(y_test, svm_preds))
```


 SVM Accuracy: 0.958368932038835

SVM Report:

	precision	recall	f1-score	support
0	0.96	0.96	0.96	6793
1	0.96	0.95	0.96	6082
accuracy			0.96	12875
macro avg	0.96	0.96	0.96	12875
weighted avg	0.96	0.96	0.96	12875

```
print("KNN Accuracy:", accuracy_score(y_test, knn_preds))
```

```
print("KNN Report:\n", classification_report(y_test, knn_preds))
```

 KNN Accuracy: 0.930873786407767

KNN Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.93	0.94	0.93	6793
1	0.93	0.93	0.93	6082